




MLOps

**K Kiran Kumar–
INSOF
Senior Assistant Professor**



<https://stanford-cs329s.github.io/>

What is a Productionized model

A **productionized model** is a **prediction service** that makes predictions on new (unseen) data. It should support **automated** training, testing, and deployment. It should be scalable, available, and if the model is online it should be able to enrich feature vectors using a feature store.




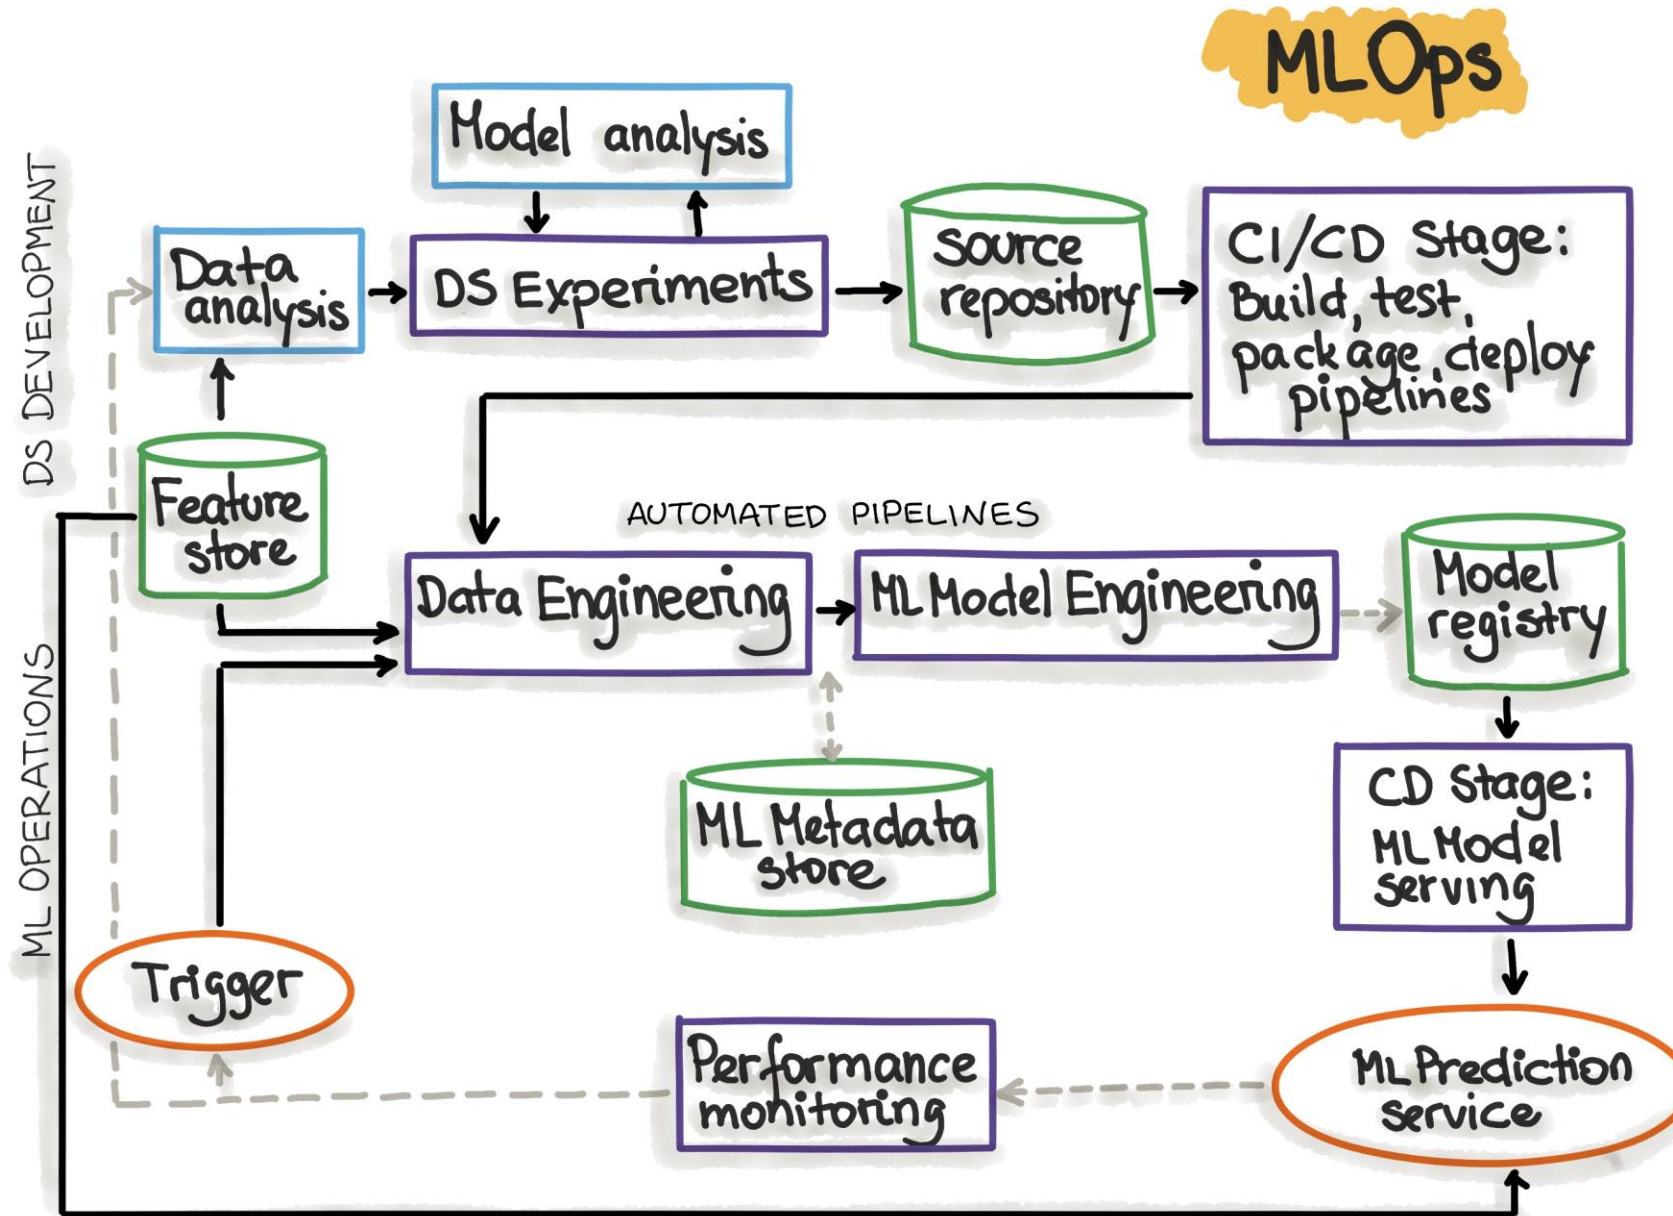
If the model is a bus - the “bus service” is “prediction service”.

What is MLOps

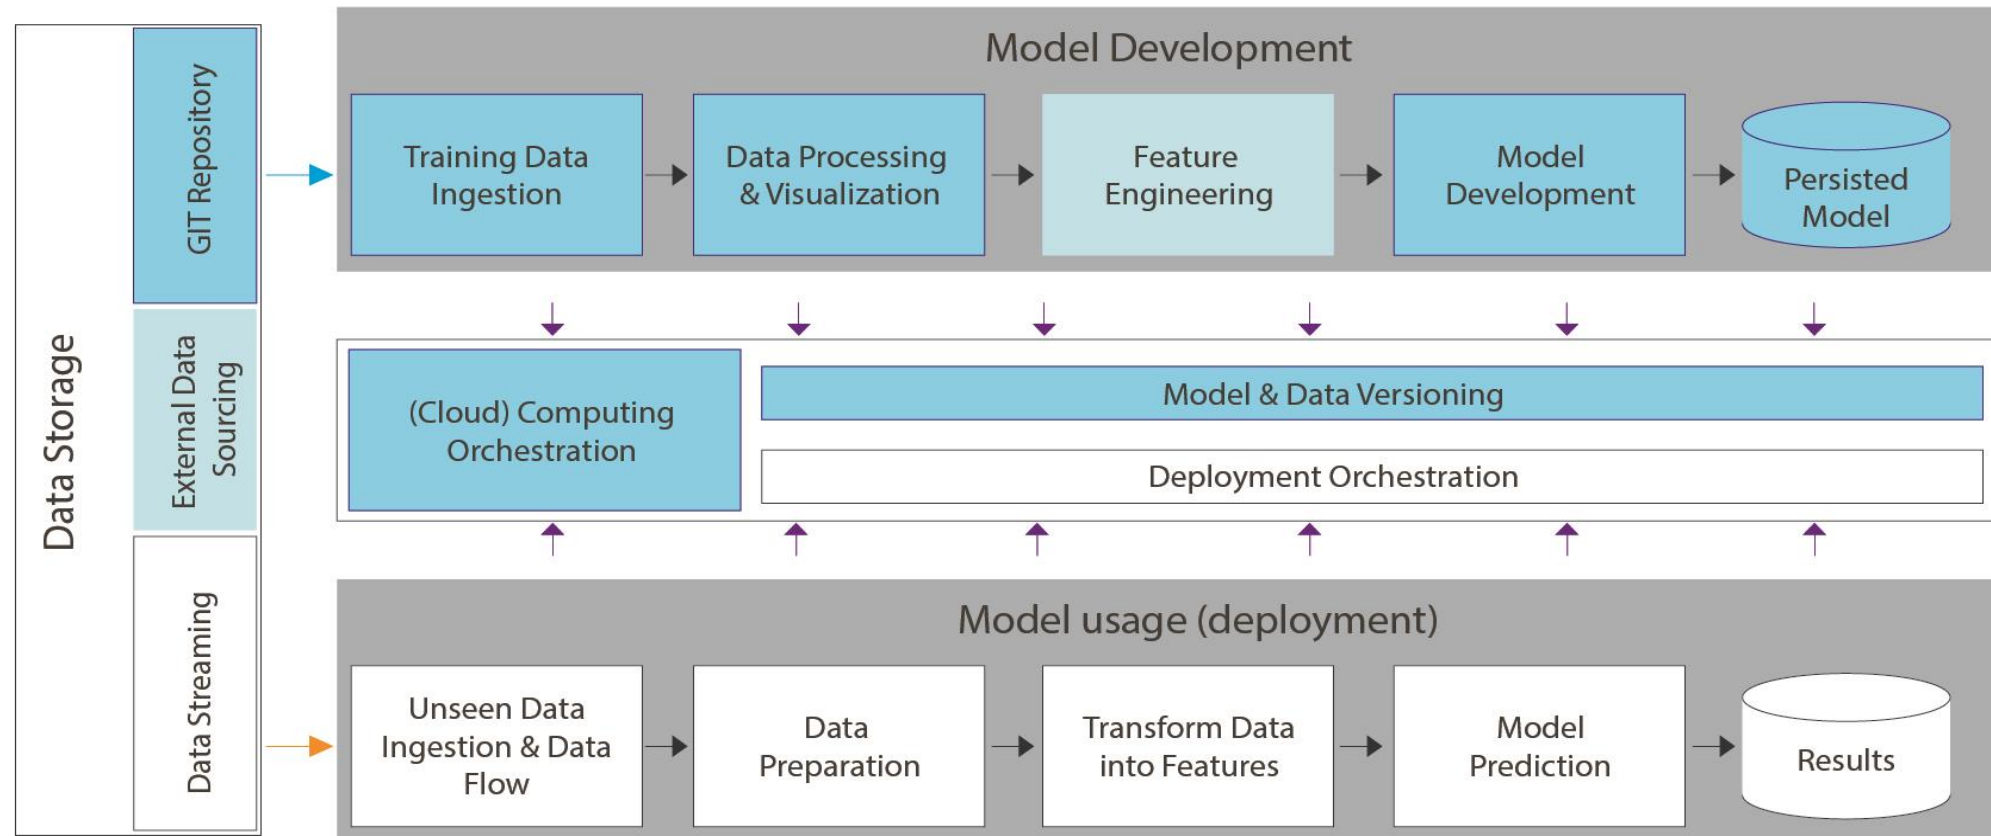
The Continuous Delivery Foundation's SIG-MLOps defines MLOps as “the extension of the DevOps methodology to include Machine Learning and Data Science assets as first class citizens within the DevOps ecology”.

MLOps is not the chat, it's the chat Bhandar





Model development & usage process

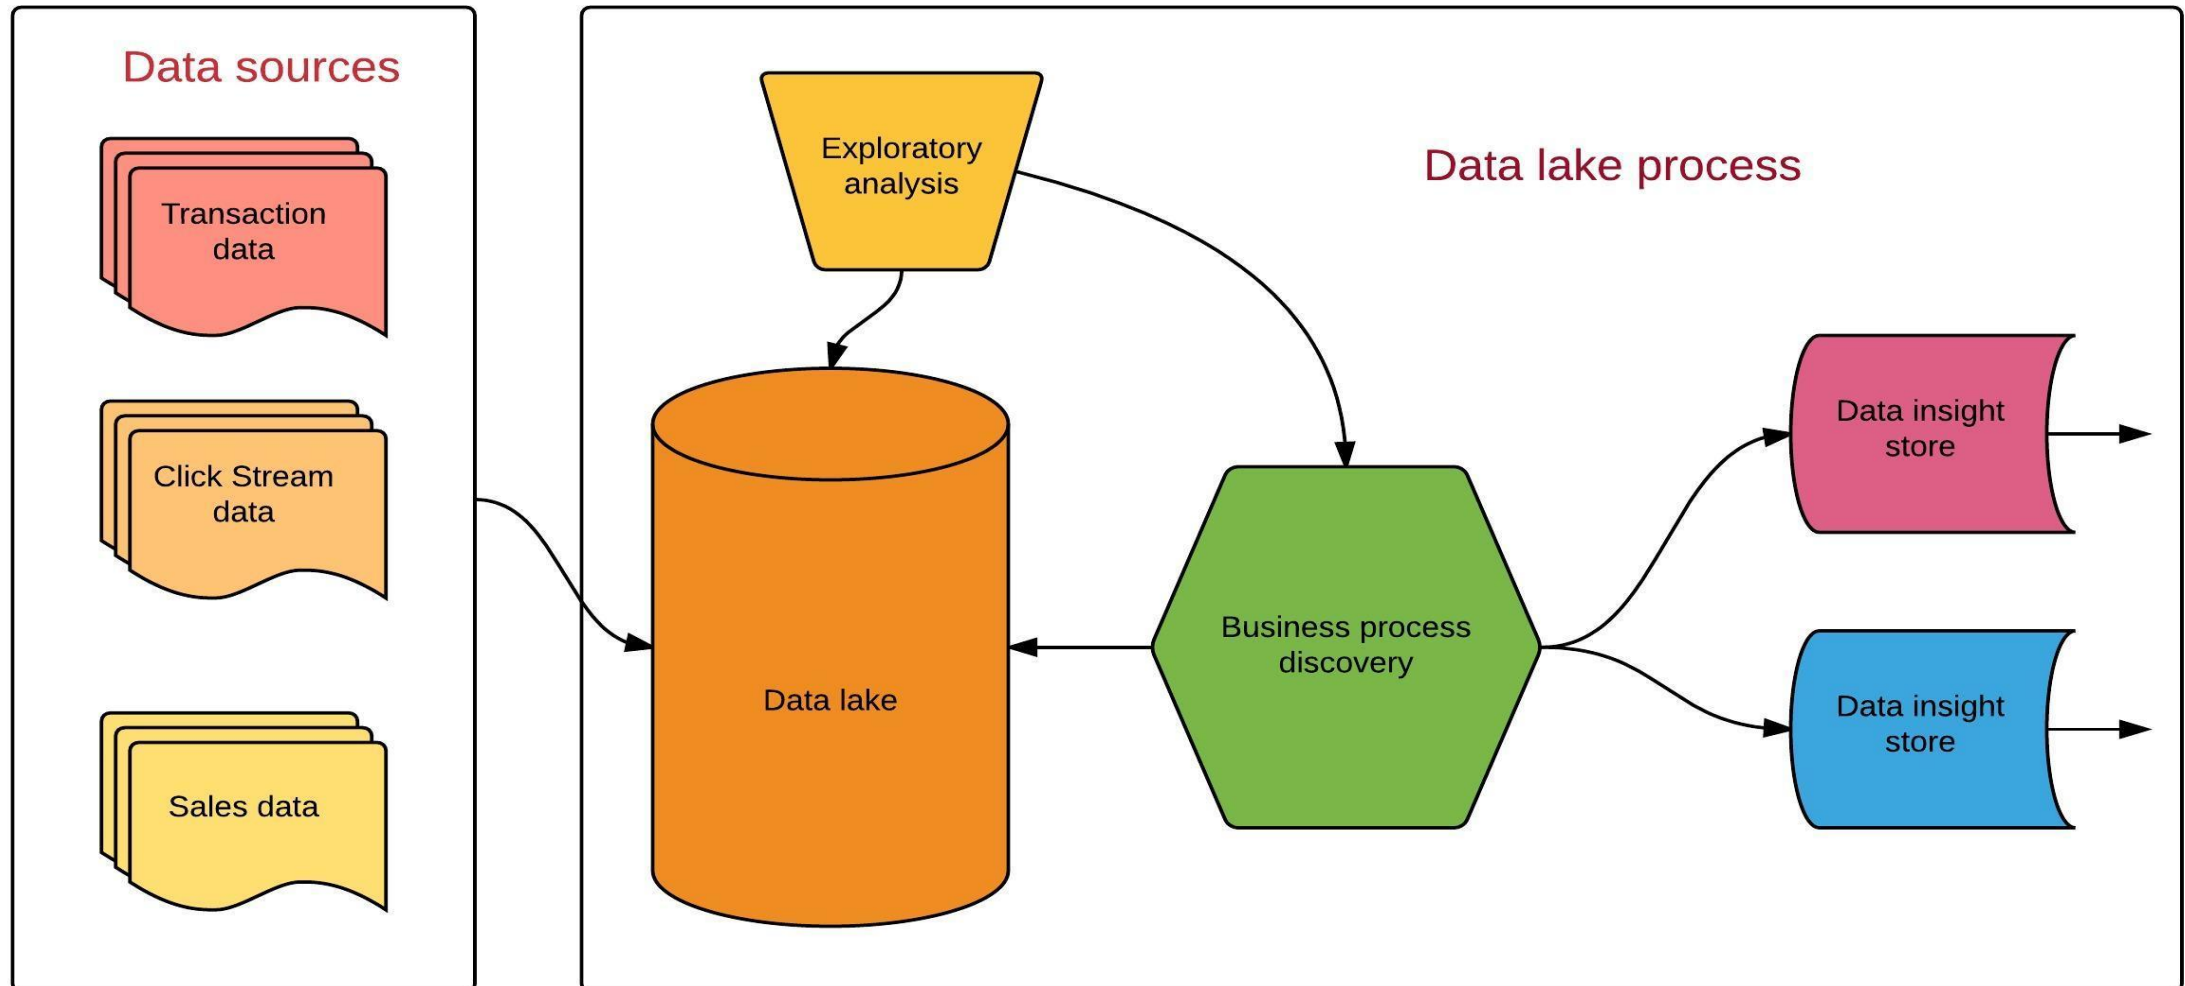



The Need for MLOps

- Data is constantly changing
 - Need to constantly update training data and models
- Multiple teams need to coordinate to put models in production
 - Data Engineers -> Data Scientists -> ML Engineers
- Infrastructure support needed
 - Feature engineering -> model training -> model serving
 - Orchestration support for pipelines
 - Data and code versioning

Project Setup

- ❖ Goals
- ❖ User Experience
- ❖ Performance constraints
- ❖ Evaluation
- ❖ Personalization
- ❖ Project Constraints





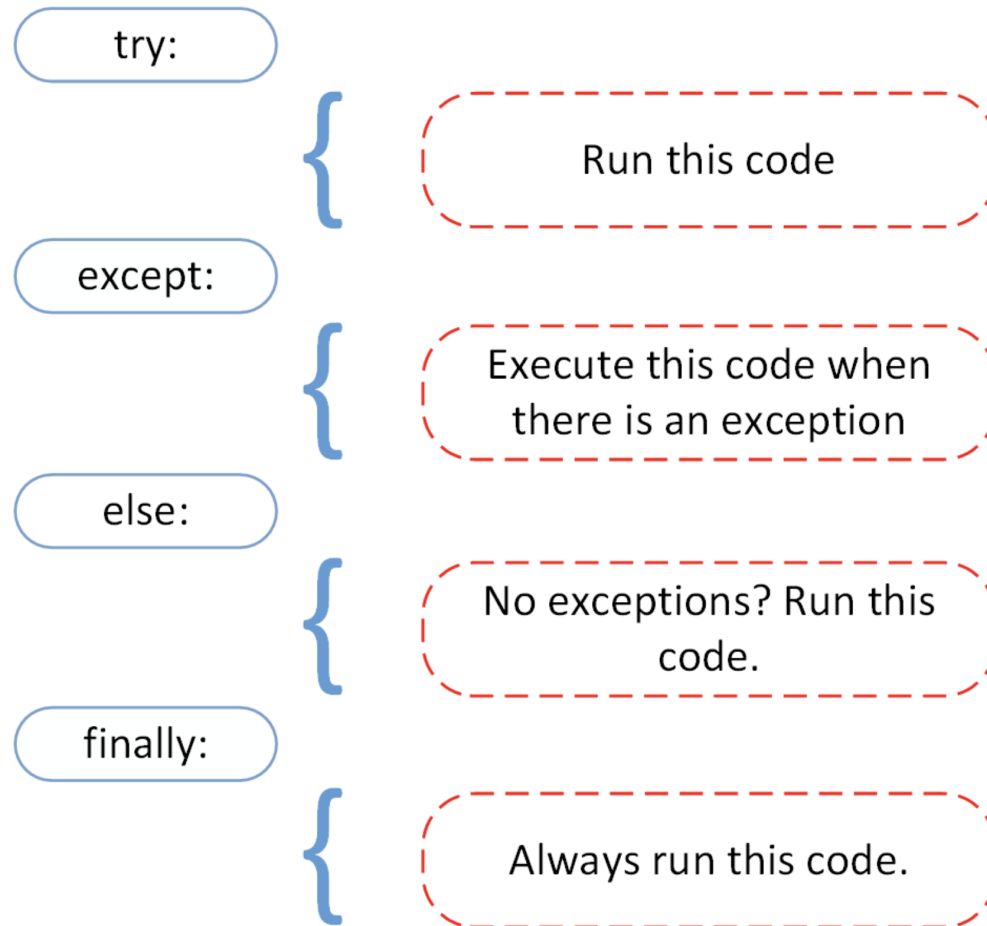
Feature Store


The feature store is the data warehouse for Data Science - it is a central vault for storing documented, curated, and access-controlled features that can be used across many different models. The feature store ingests data from the Enterprise's many different sources after transforming, aggregating, and validating the data.

Feature Store

	Used by	What is updated, when and how?
<i>Training Data</i>	Model Training	Generate new train/test data when new labelled features become available in the feature store.
<i>Analytical Models</i>	Batch Apps	Run batch jobs on-demand or periodically when new unlabelled features become available in the feature store.
<i>Operational Models</i>	Online Apps	Periodically update the online features in the feature store, and query those features to build feature vectors when online predictions are needed.







Machine Learning Model Management

Features	DVC	mlFlow	Sacred
<i>Language Support</i>	Universal using Shell	Python, R, Java , REST	Python
<i>Control</i>	DAG using scripts	Python package	Annotations
<i>Storage Backend</i>	File Storage	File Storage	MongoDB + others
<i>License</i>	Apache 2	Apache 2	MIT
<i>Notebook support</i>	Yes	Yes	Only Dry-run
<i>Model Storage</i>	Yes	As Artefacts	Only Artefacts
<i>Input Data Storage</i>	Yes	Only Artefacts	Yes
<i>Parameter Tracking</i>	Yes	Yes	Yes



Open source machine learning platform

- Works with any ML library, algorithm, language, etc
- *Open interface* design (use with any code you already have)

mlflow Tracking

Record and query
experiments: code,
data, confs, results

mlflow Projects

Packaging format
for reproducible
runs and workflows

mlflow Models

General format
that standardizes
deployment paths

new

mlflow Model Registry

Centralized model
management,
review & sharing

THANK YOU